BUILDING MULTIMODAL SEARCH AND RAG

- MULTIMEDIA CONTENT IS ALL AROUND US;
- DATA FROM DIFFERENTS SOURCES;
- TRAINING MULTIMODAL MODELS: START WITH SPECIALIST MODELS;
  - (a) TEXT ENCODER;
  - (b) IMAGE ENCODER;
  - (c) AUDIO ENCODER;
  - (d) VIDEO ENCODER;
- SIMILAR CONCEPTS = SIMILAR VECTORS;

  (DATA) (VECTORS REPRESENTATIONS)

  ↳ UNIFY THE SPECIALIST MODELS;
- UNIFY THE MODELS USING CONTRASTIVE LEARNING:

  ↳ PROCESS TO TRAIN ANY EMBEDDING MODEL;

  ↳ UNIFY MULTIPLE MODELS;

  ↳ CREATE ONE VECTOR SPACE;

  ↳ TUNE MODELS BY PROVIDING CONTRASTIVE EXAMPLES;

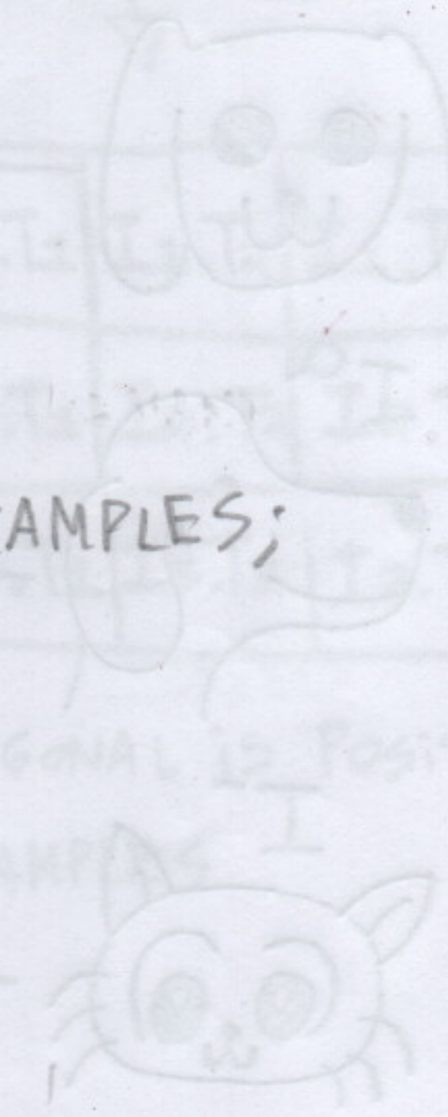

- E.g: ANCHOR → "HE COULD SMELL THE ROSES"

  POSITIVE EXAMPLE → "A FIELD OF FRAGRANT FLOWERS"

  NEGATIVE EXAMPLE → "THE LION ROARED MAJESTICALLY"

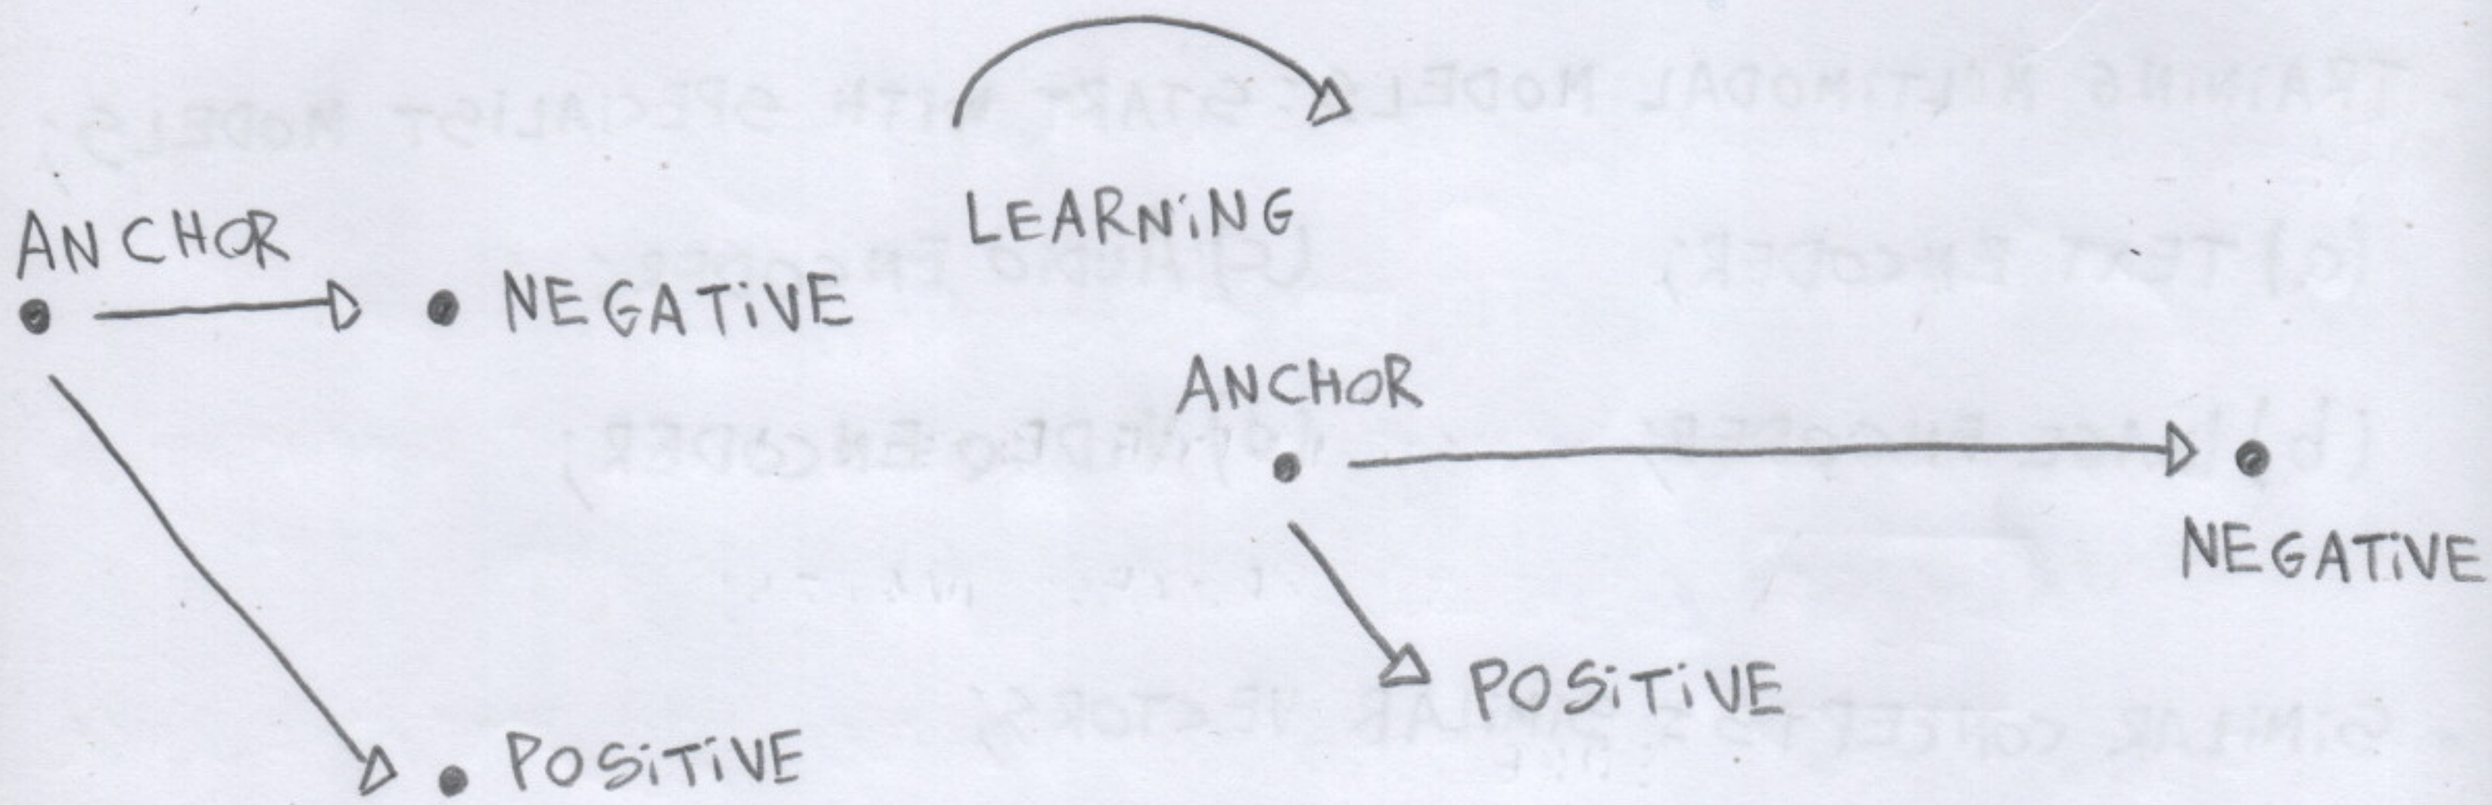

↳ SAME EXAMPLE WITH IMAGES, AUDIO AND VIDEO;

↳ PUSH NEGATIVE EXAMPLE;

↳ PULL POSITIVE EXAMPLE;

↳ CONTRASTIVE LOSS FUNCTION

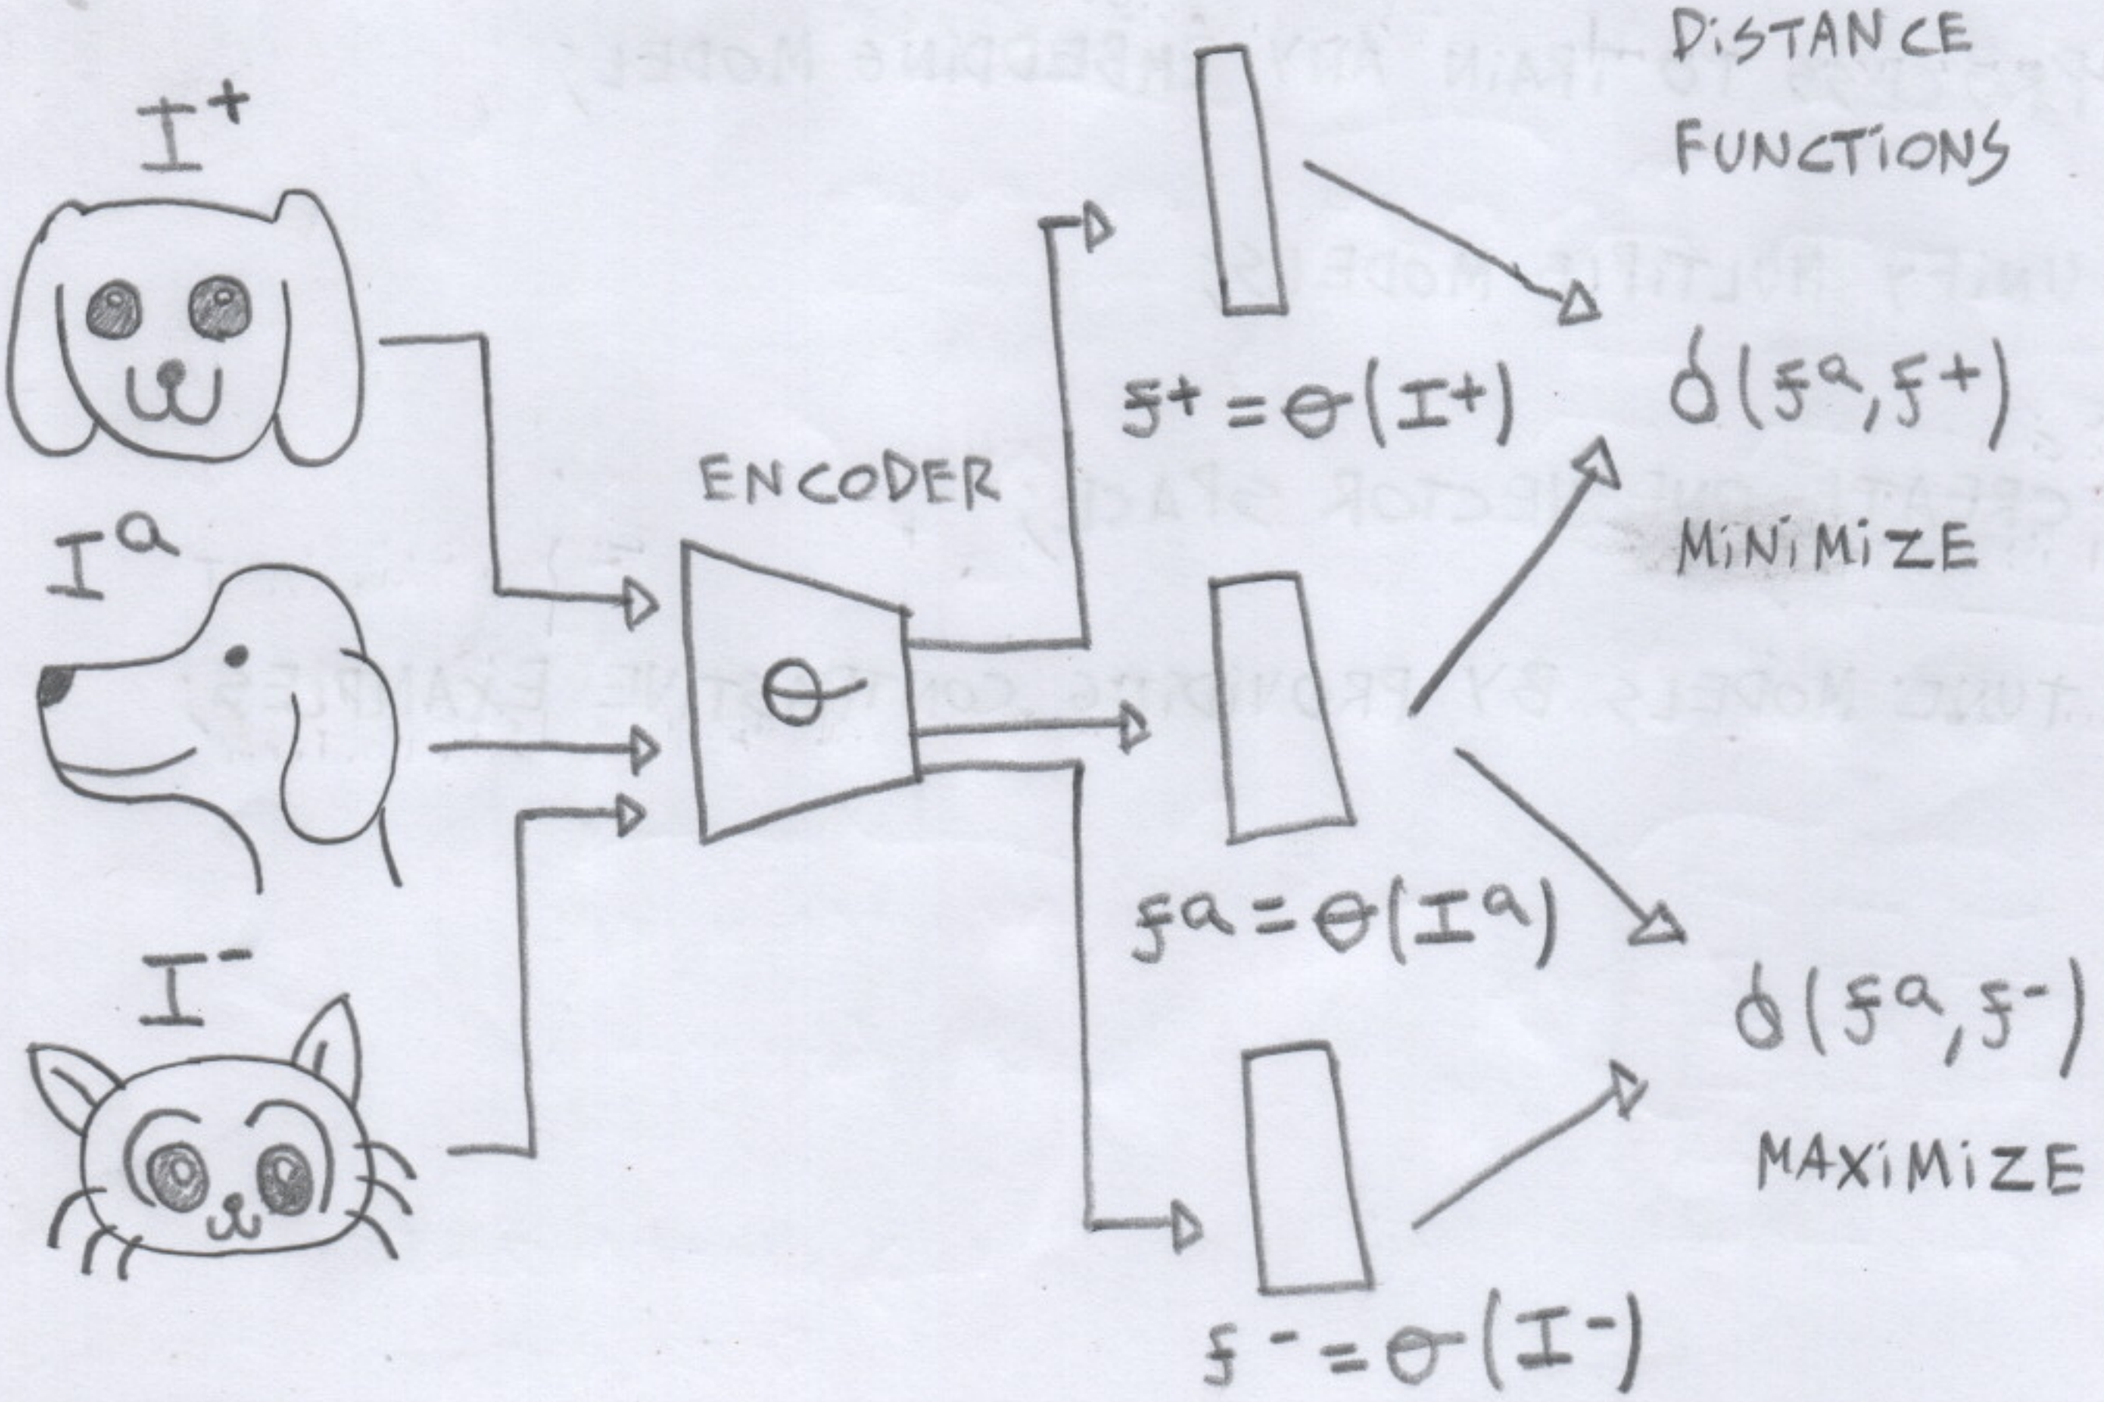

E.g: APPLY CONTRASTIVE LOSS IN TEXT AND IMAGE MULTIMODAL DATA.

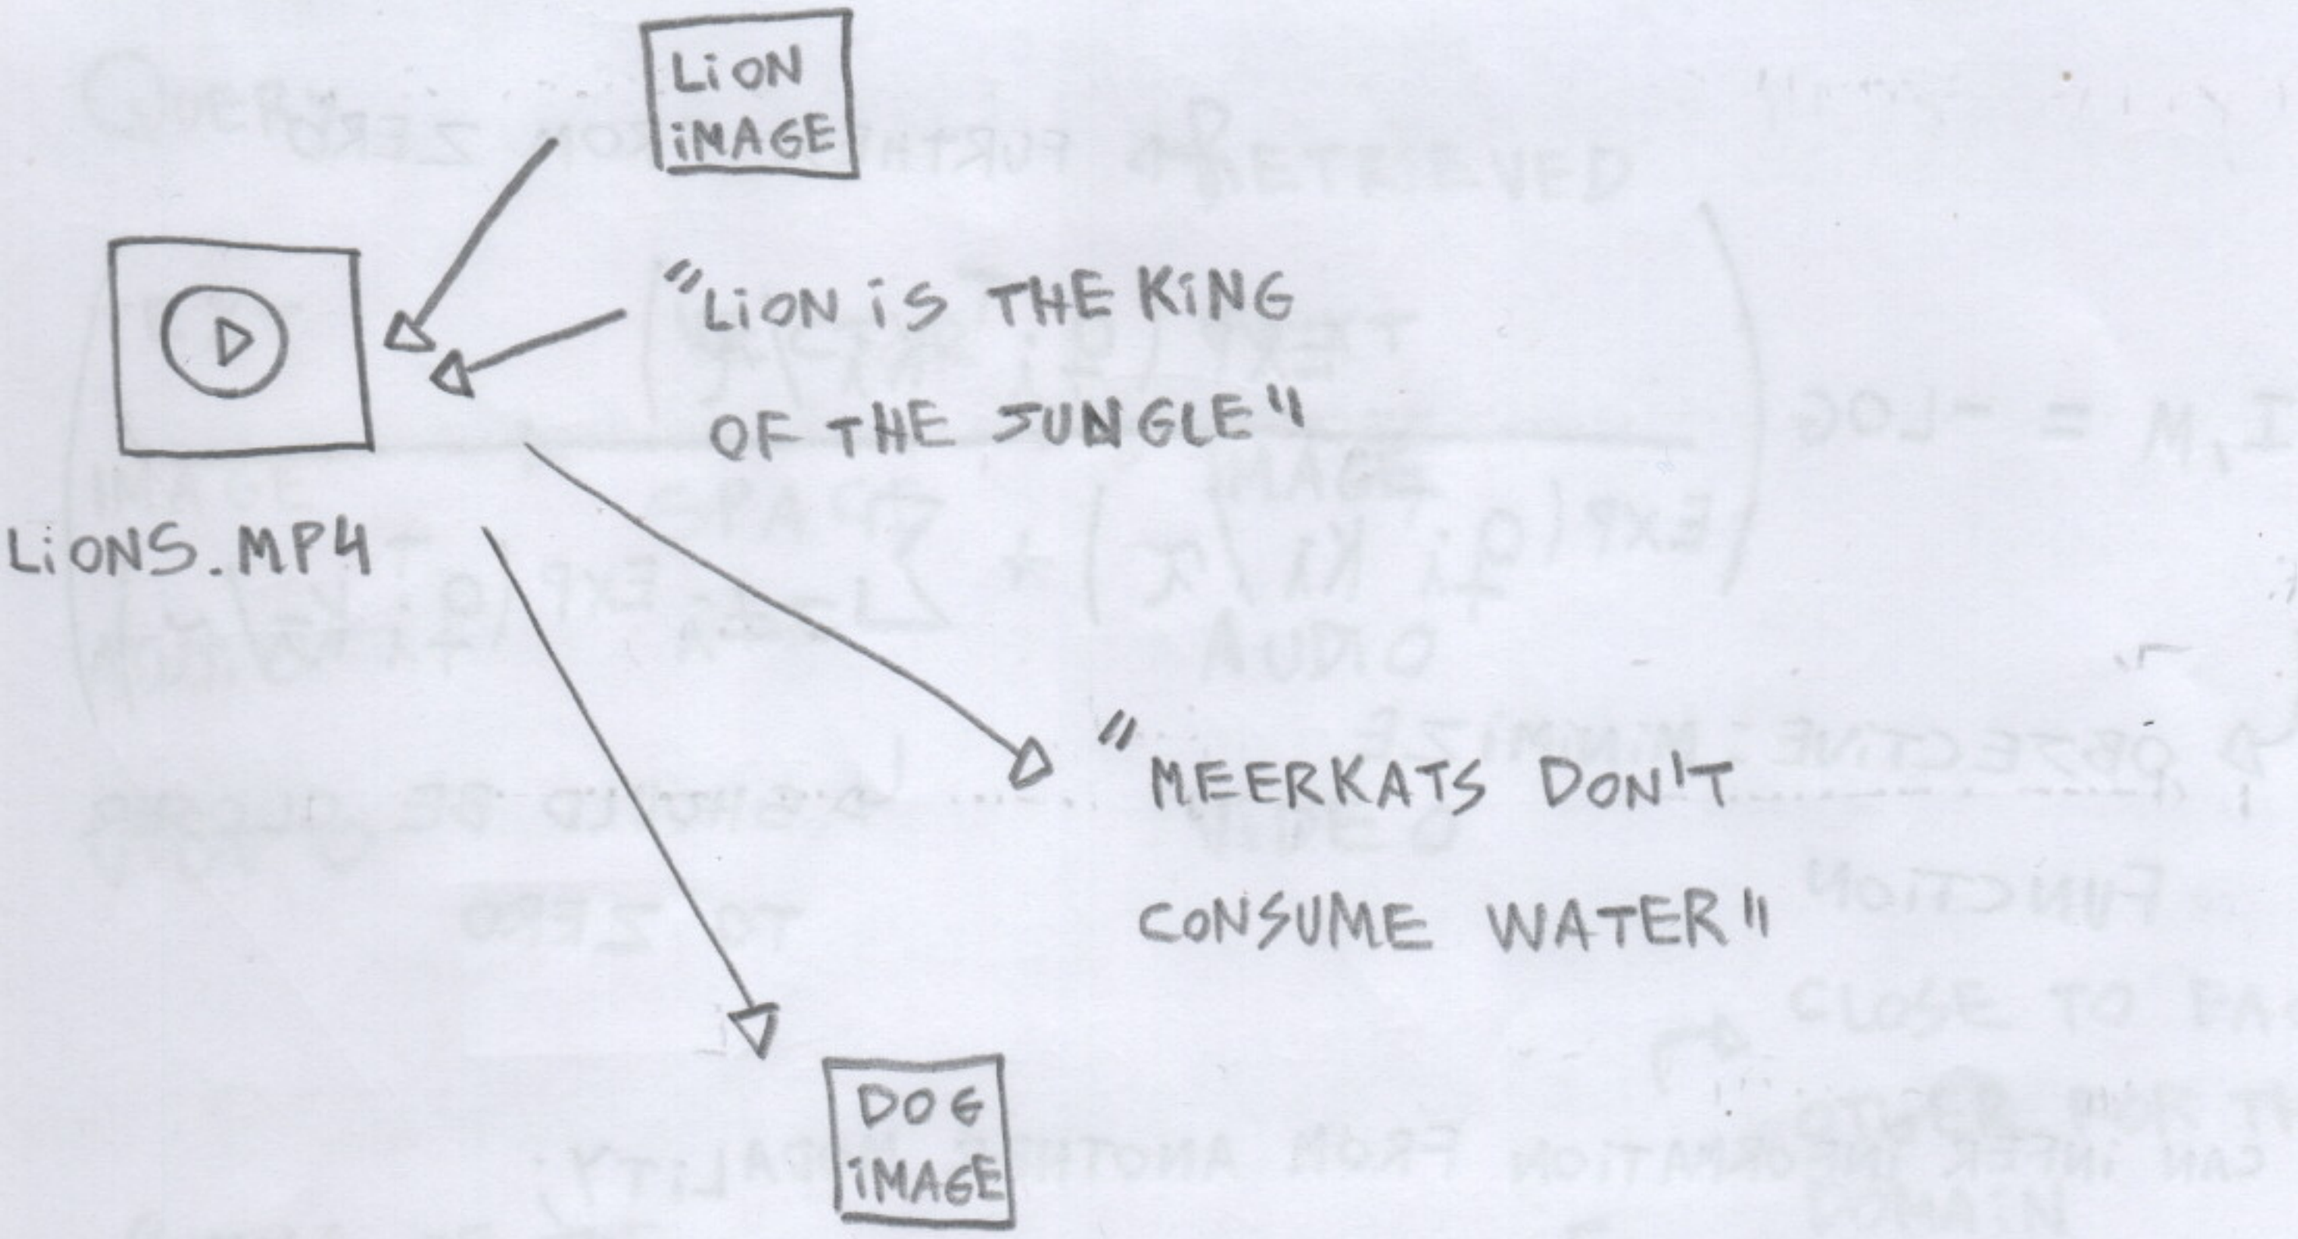

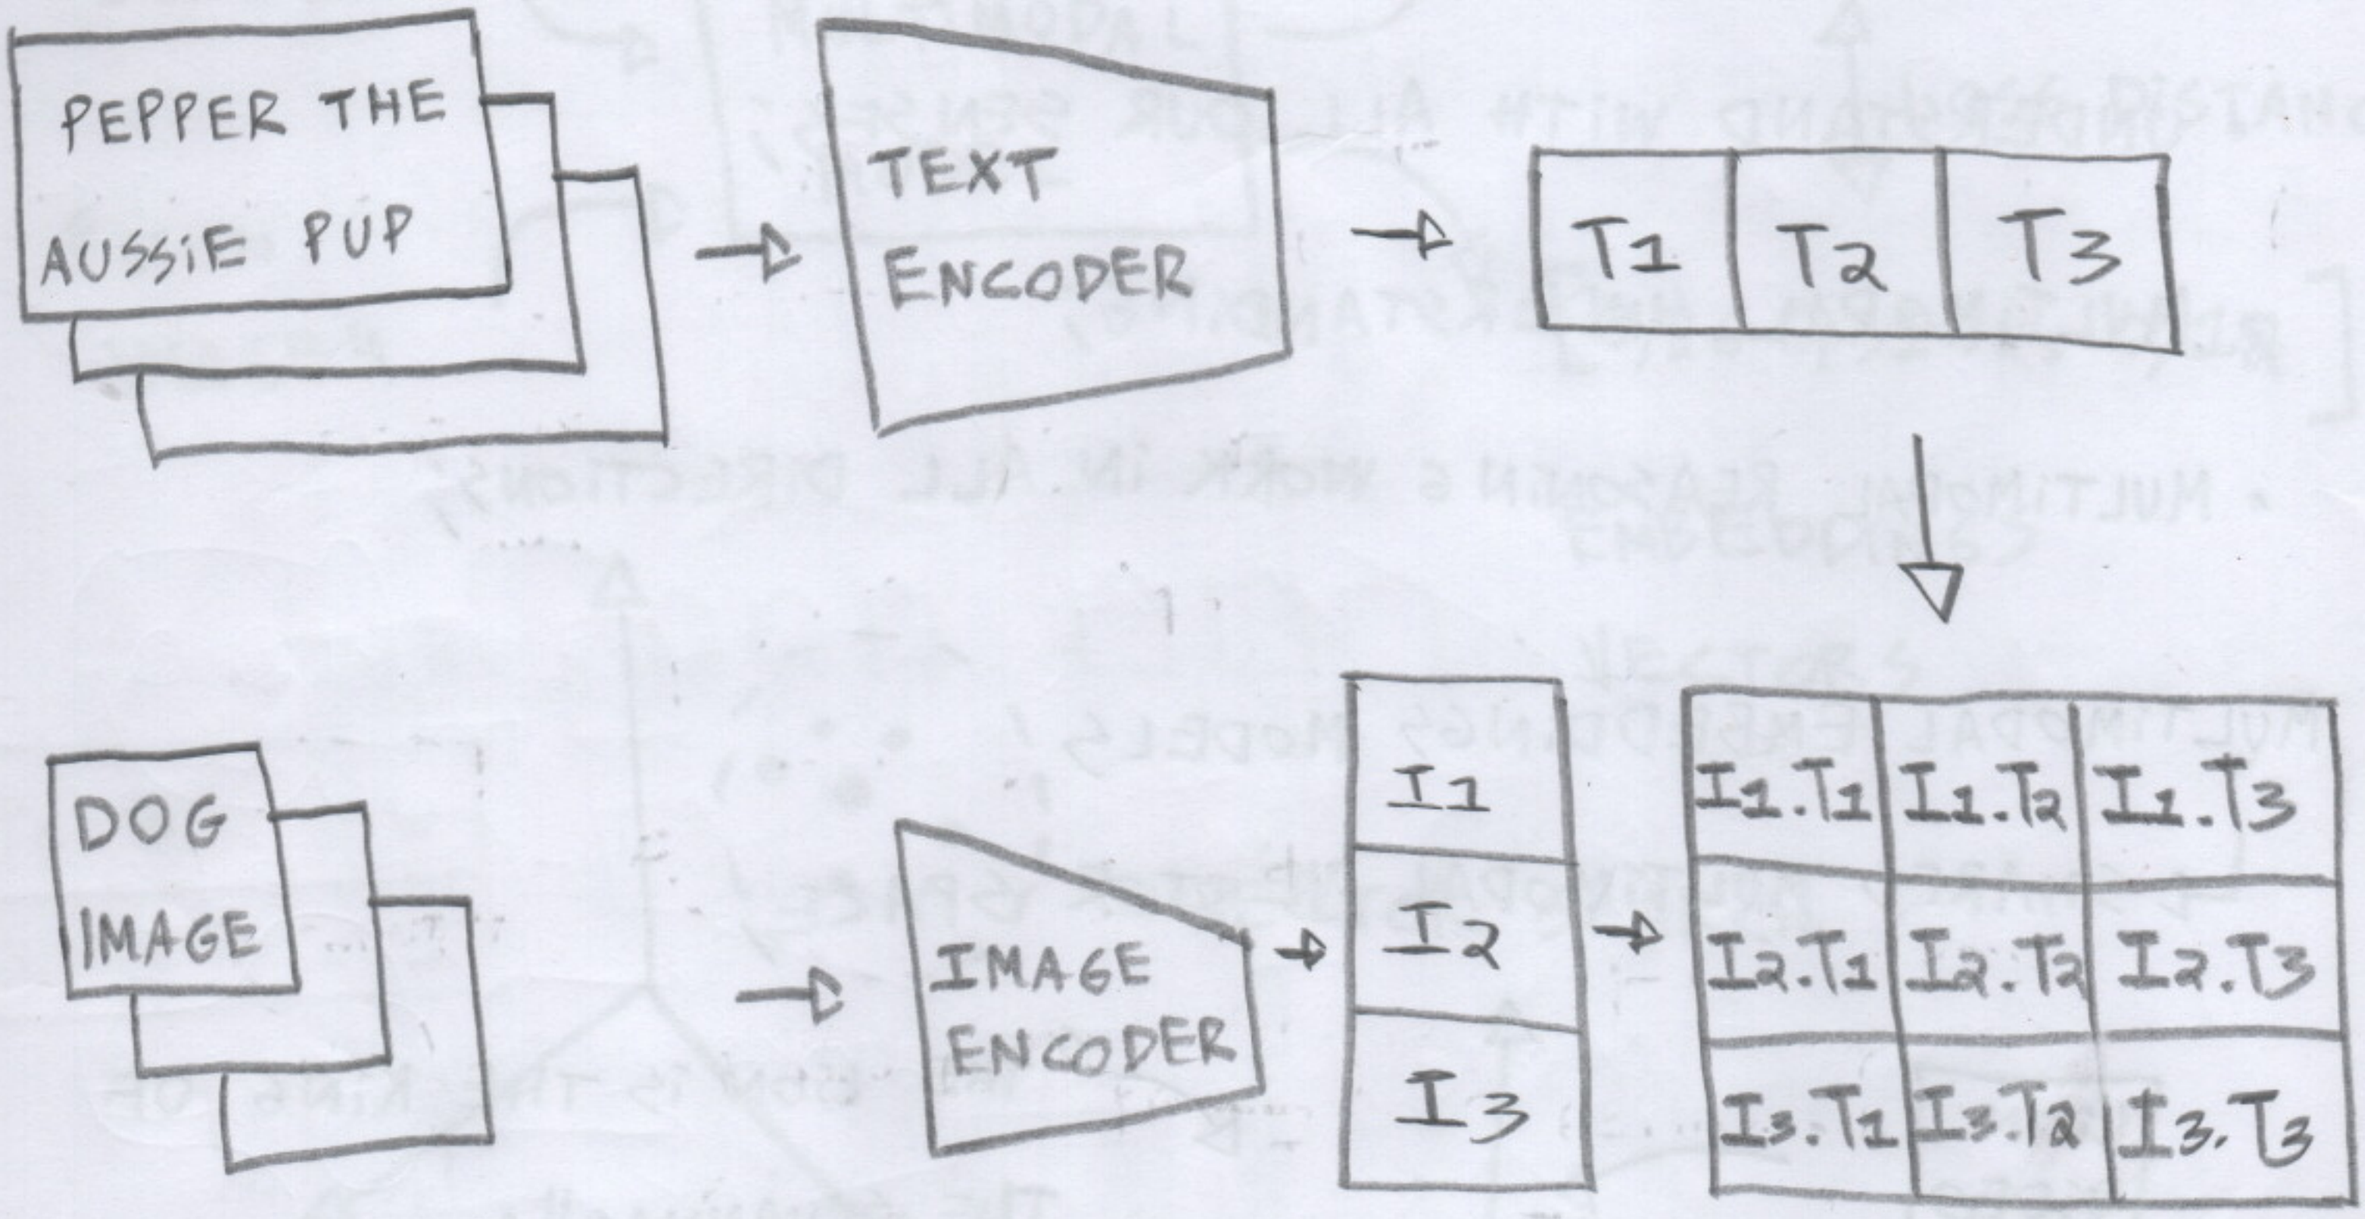

DIAGONAL IS POSITIVES EXAMPLES

$q_i = f(\text{"IMAGE OF LION"})$ AND $k_i = g(\text{"VIDEO OF LION"})$

$$L_{I,M} = -\log\left(\frac{\exp(q_i^T k_i/\tau)}{\exp(q_i^T k_i/\tau) + \sum_{j\neq i}\exp(q_i^T k_j/\tau)}\right)$$

↳ FURTHEST FROM ZERO

↳ OBJECTIVE: MINIMIZE FUNCTION

↳ SHOULD BE CLOSER TO ZERO

- WE CAN INFER INFORMATION FROM ANOTHER MODALITY;
- MULTIMODAL REASONING AND RETRIEVAL:
  - UNDERSTAND WITH ALL OUR SENSES;
  - MULTIMODAL UNDERSTANDING;
  - MULTIMODAL REASONING WORK IN ALL DIRECTIONS;
- MULTIMODAL EMBEDDINGS MODELS

  ↳ SHARED MULTIMODAL VECTOR SPACE;

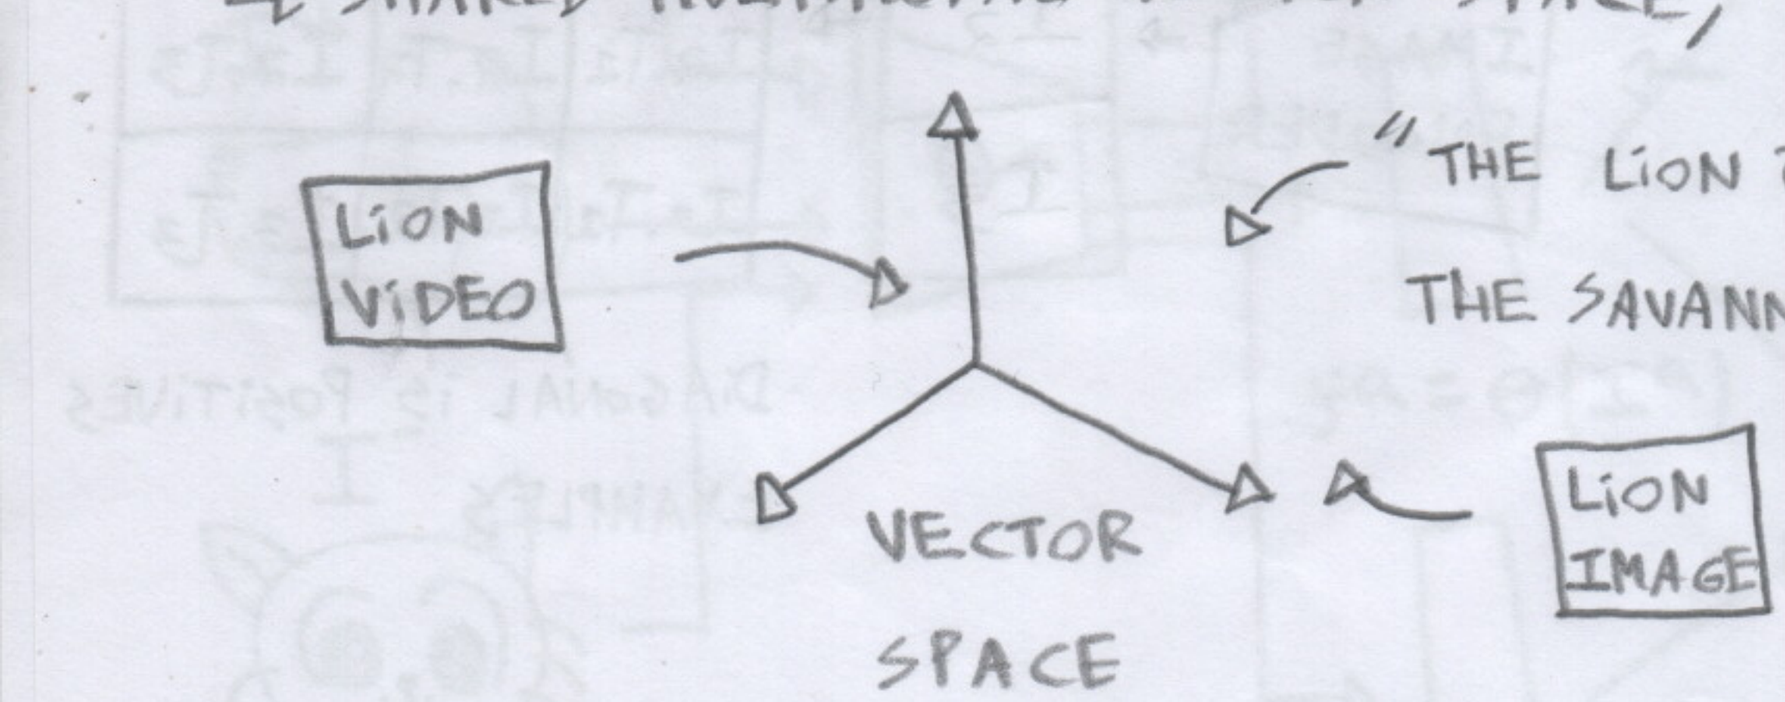

ANY TO ANY SEARCH

QUERY

TEXT
IMAGE
AUDIO
VIDEO

→ VECTOR SPACE →

RETRIEVED

TEXT
IMAGE
AUDIO
VIDEO

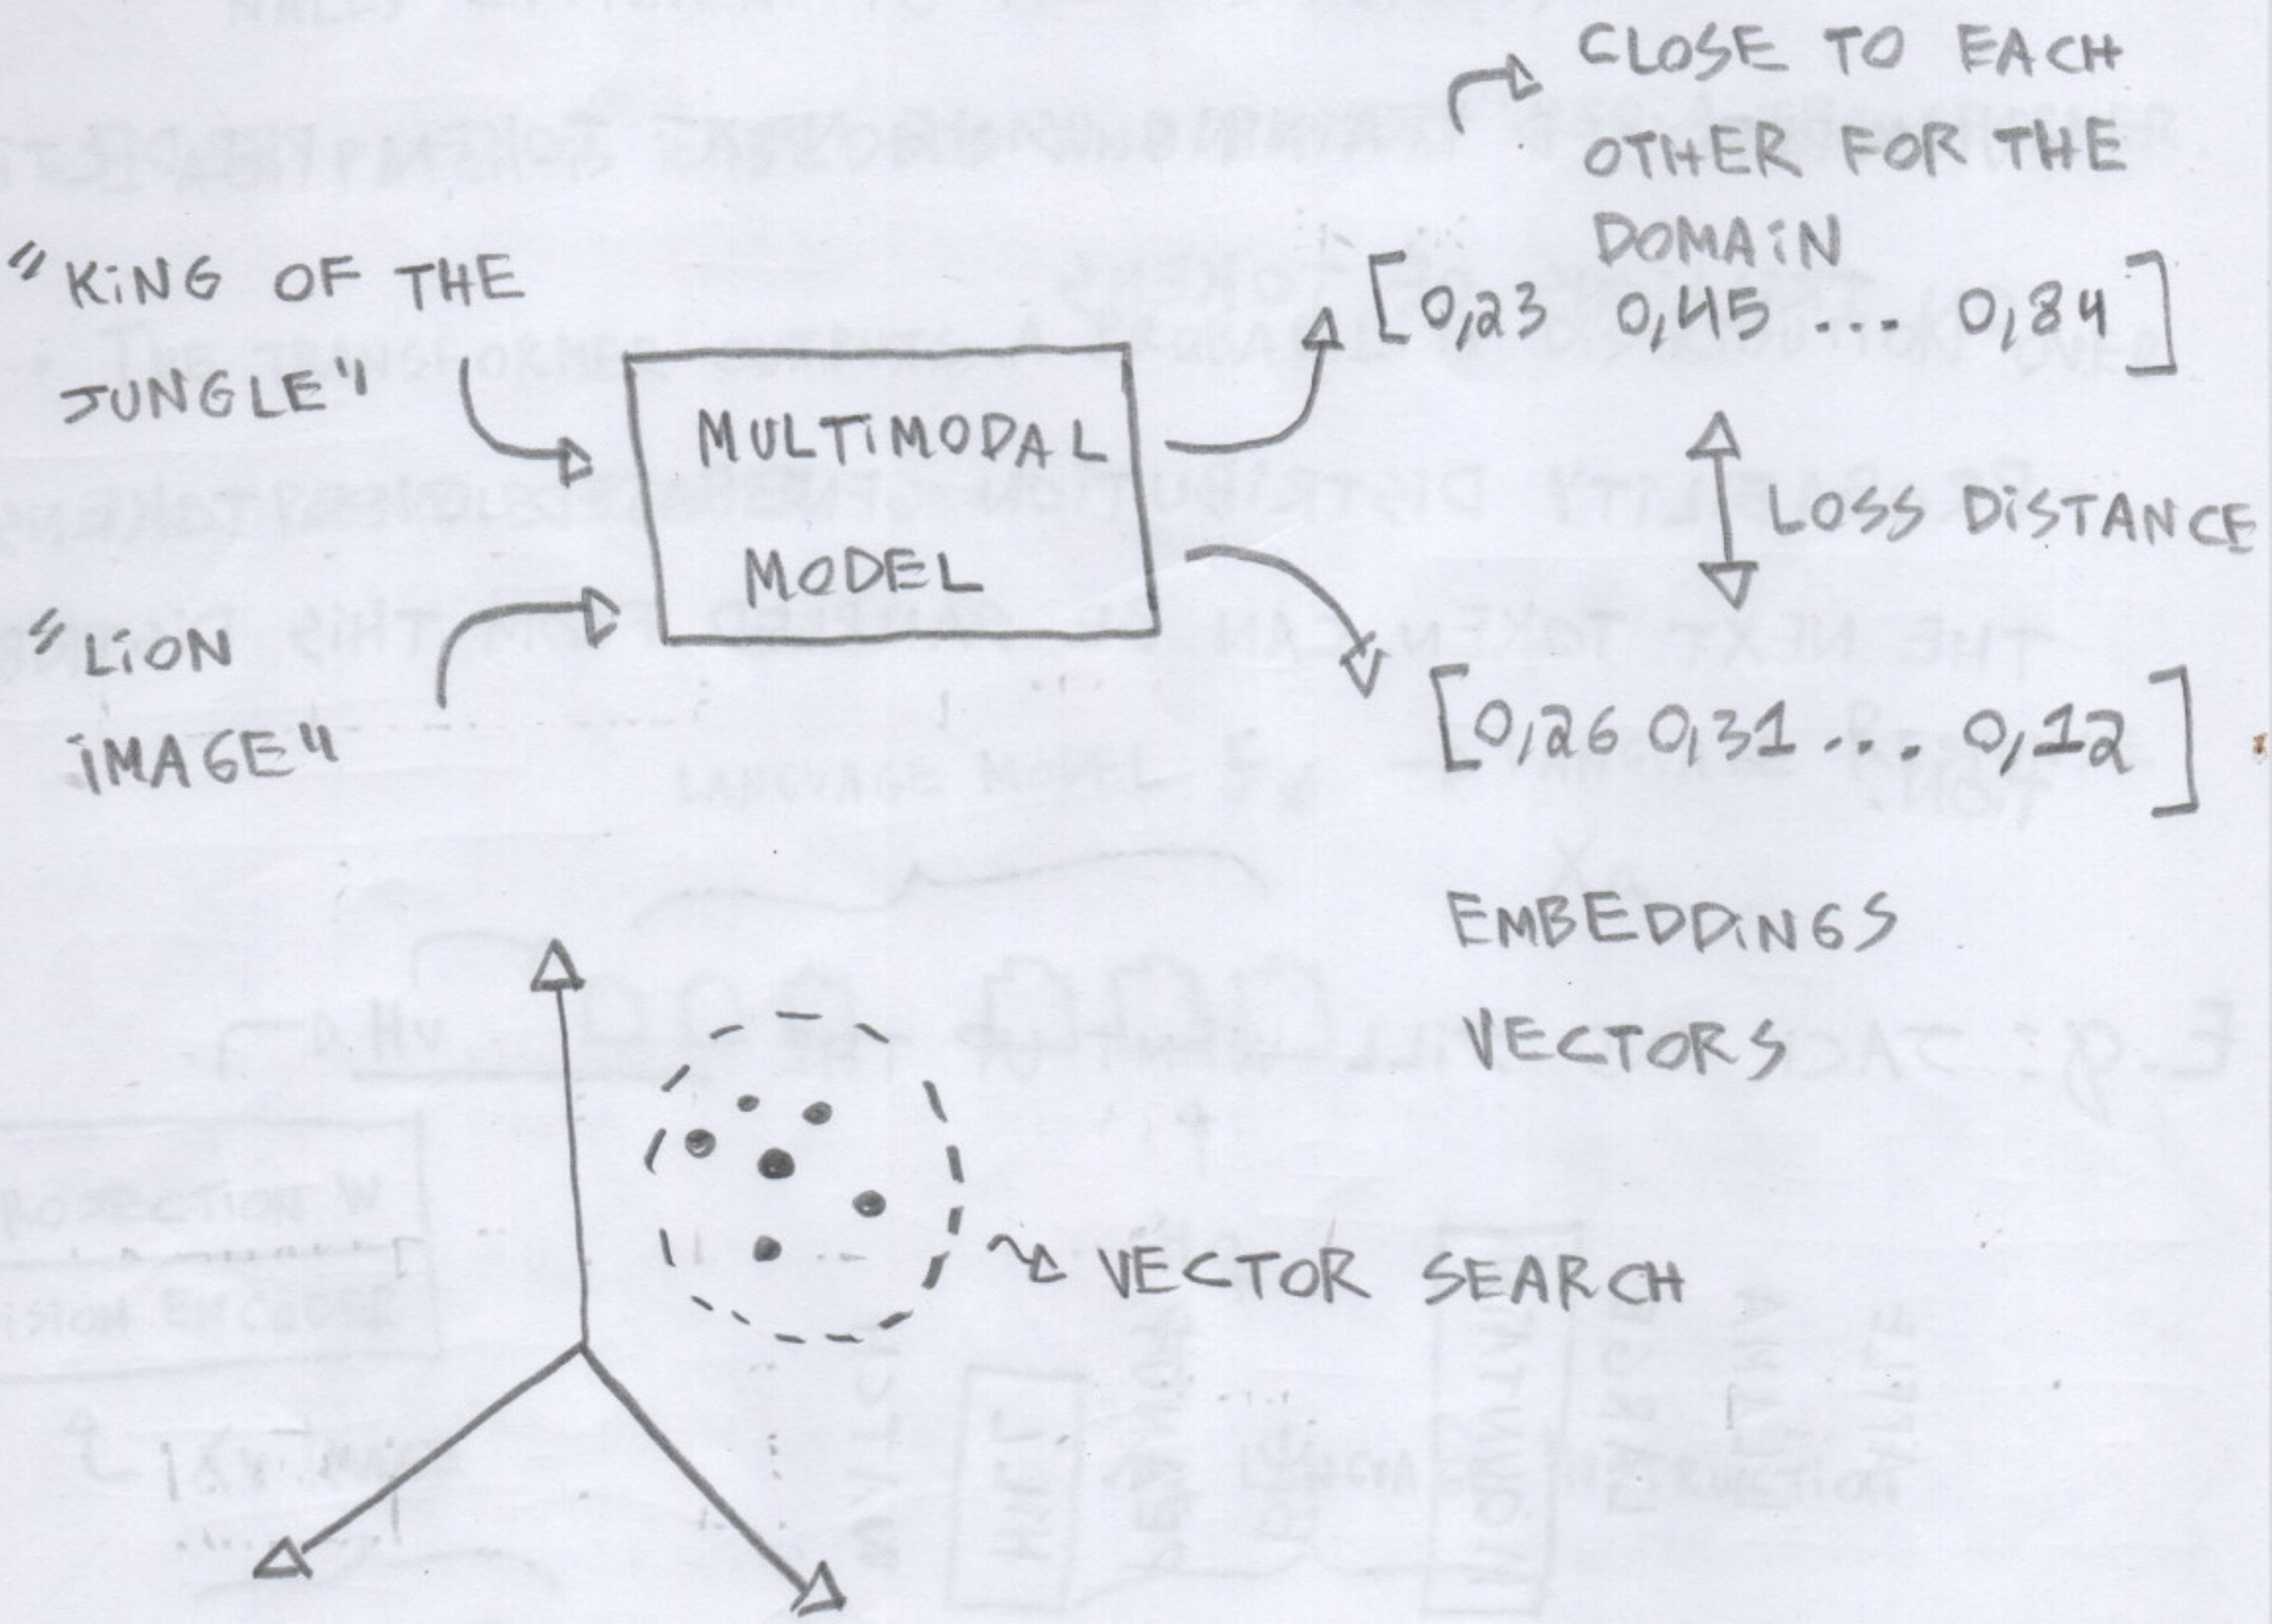

- HOW DO LARGE LANGUAGES MODELS WORK?

  ↳ THE MAJORITY OF CURRENT LLMS ARE GENERATIVE PRE-TRAINED TRANSFORMERS (GPT).

  - AUTOREGRESSIVE: THEY GENERATE TEXT ONE TOKEN AT A TIME;
  - FUTURE TOKENS ARE CONDITIONED ONLY ON PREVIOUSLY PROVIDED OR GENERATED TOKENS;
  - UNSUPERVISED TRAINING USING NEXT TOKEN PREDICTION ON TRILLIONS OF TOKENS.
  - PROBABILITY DISTRIBUTION GENERATED OVER TOKENS: THE NEXT TOKEN CAN BE SAMPLED FROM THIS DISTRIBUTION.

E.g: JACK AND JILL WENT UP THE _____

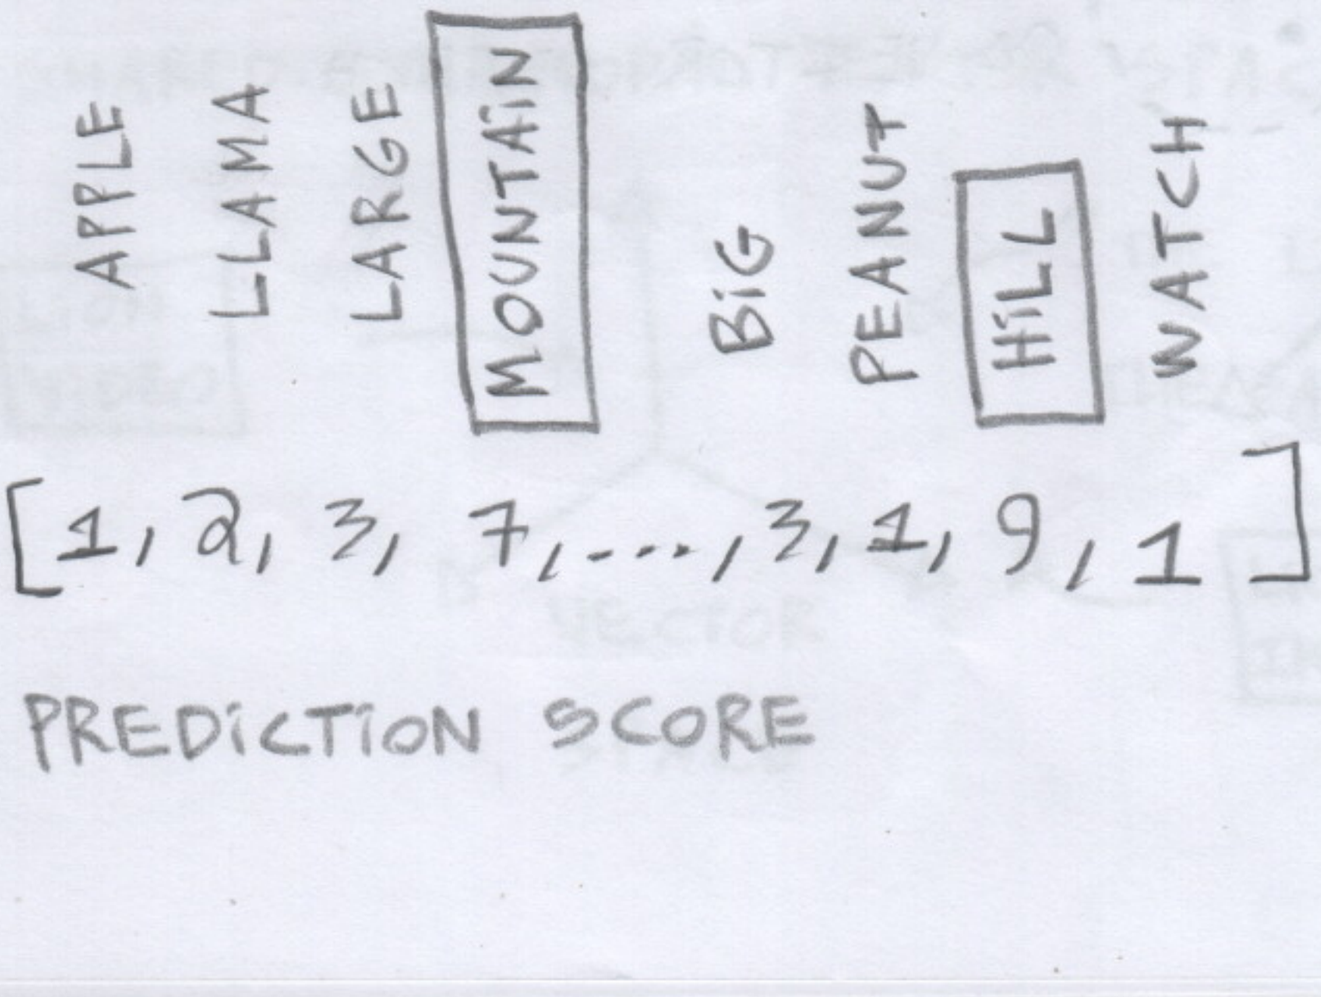

- HOW DO GPT MODELS WORK?

PROMPT: THE ROCK
- THE → EMBEDDING LOOKUP
- ROCK → TOKEN PROBABILITY

- VISION TRANSFORMERS

- IMAGES ARE CUT UP INTO PATCHES.
  - USING PATCHES INSTEAD OF PIXELS MAKE IT COMPUTATIONALLY EFFICIENT TO PROCESS IMAGES;

- EACH PATCH IS EMBEDDED AND PASSED INTO A TRANSFORMER.

- THE TRANSFORMER OUTPUTS A PROBABILITY DISTRIBUTION OVER THE POSSIBLE CLASSES.

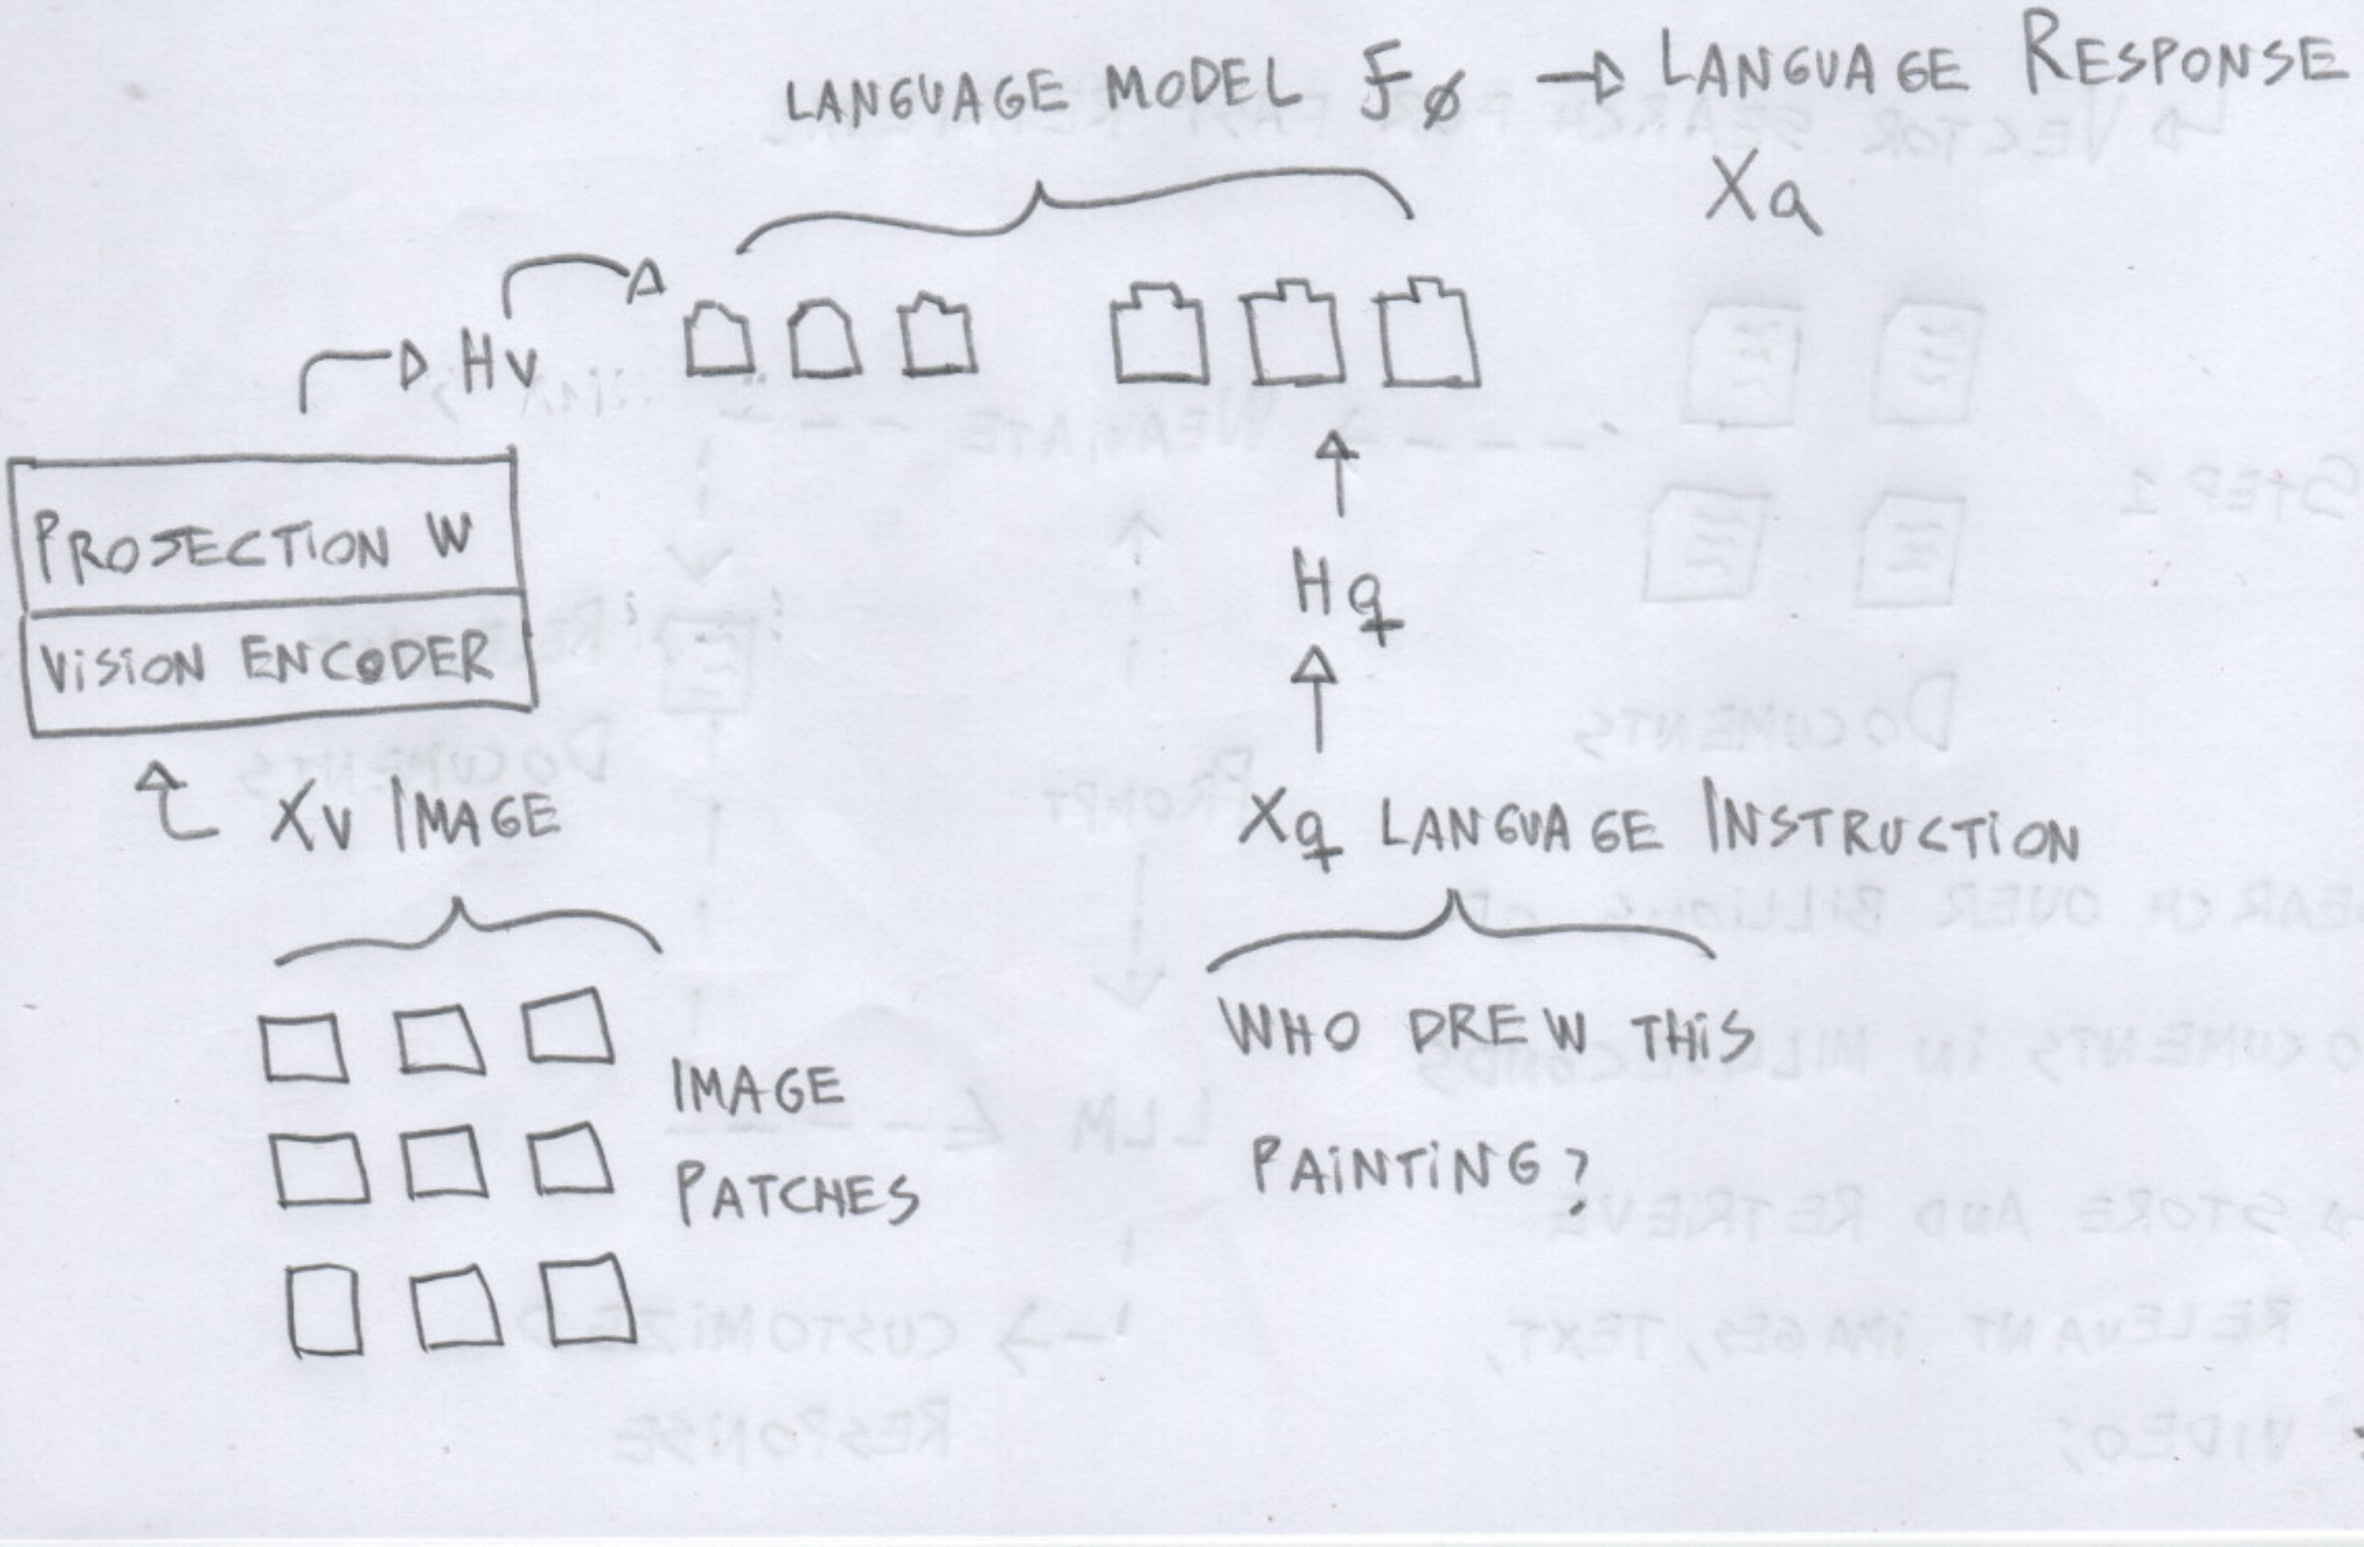

- THE PROBLEM WITH LLMS

  "YOU DON'T KNOW WHAT YOU DON'T KNOW"

  ~SOCRATES

  E.G: "WHO AT THE FAMILY PICNIC IS ALLERGIC TO NUTS?"

- RETRIEVAL AUGMENTED GENERATION (RAG)

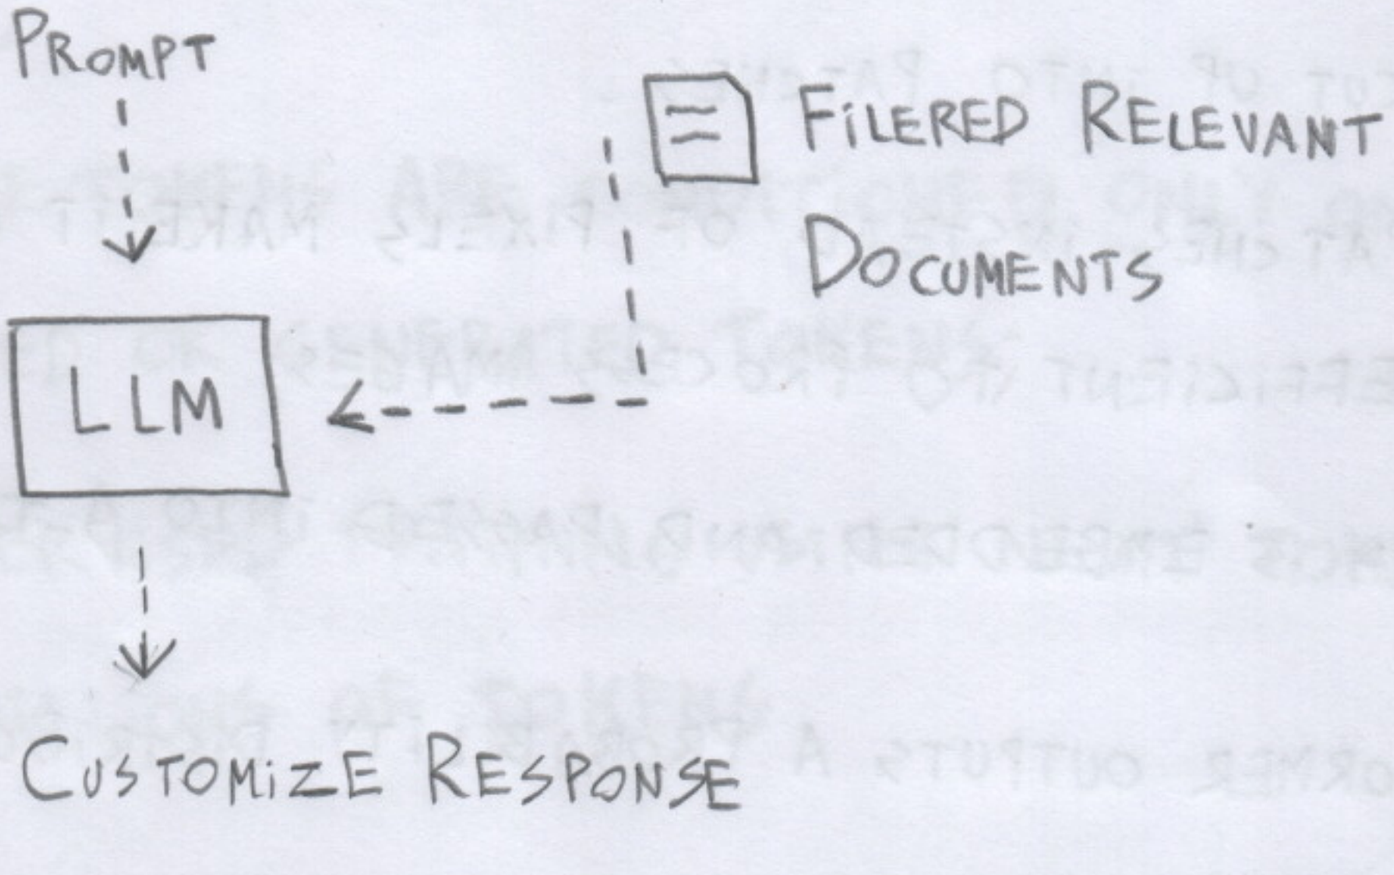

- VECTOR SEARCH FOR FAST RETRIEVAL [WITH WEAVIATE (VECTOR DB)]

  ↳ VECTOR SEARCH FOR FAST RETRIEVAL

SEARCH OVER BILLIONS OF DOCUMENTS IN MILLISECONDS

↳ STORE AND RETRIEVE RELEVANT IMAGES, TEXT, VIDEO;

- APPLICATIONS OF MULTIMODALITY IN INDUSTRY

[I] STRUCTURED DATA GENERATION

[II] TABLE CREATION

[III] UNDERSTAND LOGIC FLOWCHARTS

- MULTIMODAL RECOMMENDER SYSTEM

SEARCH IS OBJECTIVE AND RECOMMENDATION IS SUBJECTIVE.